

(Scala) Spark for Data Scientists

Ryoko Kita

Feb. 13, 2016

Cornell Tech Hackathon
NY, NY

<https://github.com/ryokokita/Spark4DS.git>

Conversation Flow


- whoami?
- End2End workflow for data scientists
- Think about your hardware to make your software GO!

Topics

- Hardware 101 for Data Scientists
 - CPU, Memory, Hard drive
 - Distributed Systems
- Software Stuff
 - Why Spark?
 - Why Scala?
- Data movement on hardware
- Airline data & data prep demo

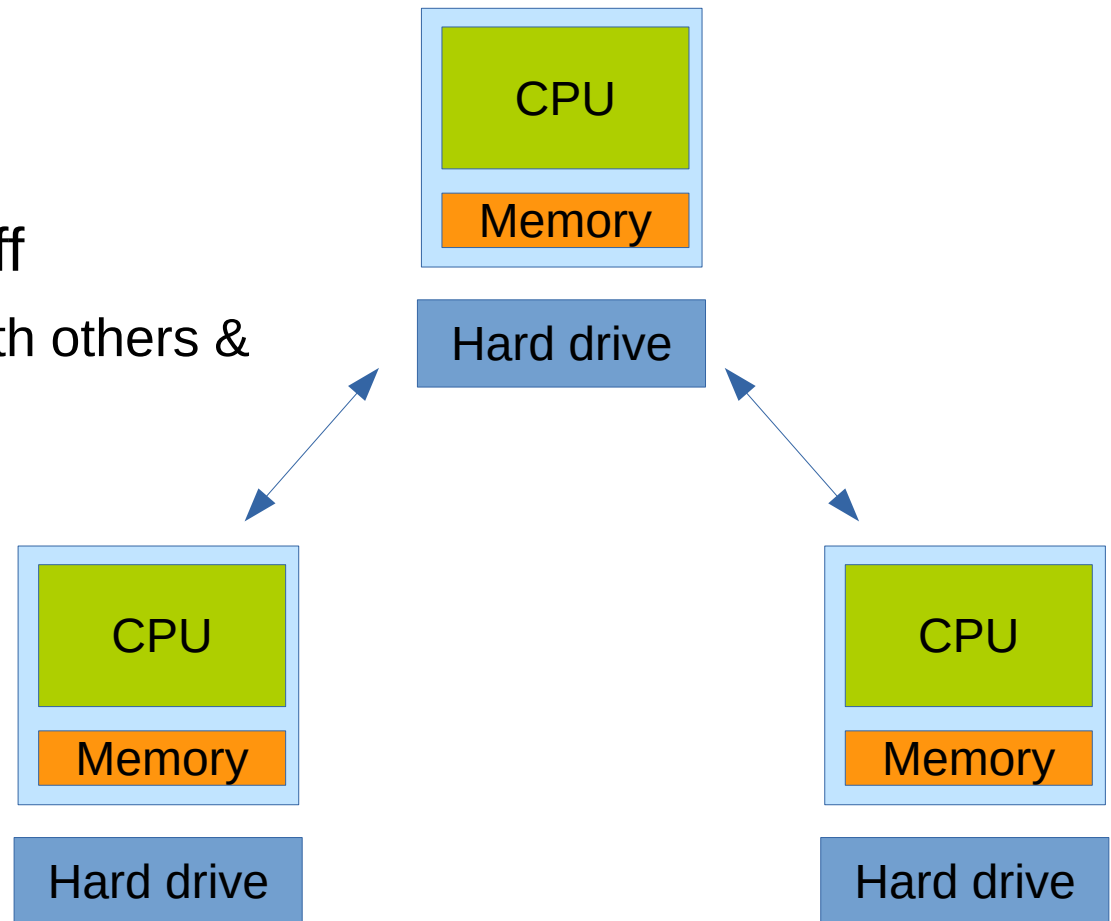
Computer Architecture 101

- CPU
 - Data transformations
 - Math-y stuff
 - Same for Spark & Hadoop MR
- Memory
 - Temporary/fast storage
 - Spark
- Hard disk/drive
 - (More) permanent storage
 - Hadoop hdfs



Distributed Systems 101

- Bunch of computers connected together
 - Via cables
 - Coordinated with software
- Generally running other stuff
 - i.e. you app has to share (with others & with other apps)
- Distributed data



Why Spark?

- One script, run anywhere:
 - Laptop, Hadoop cluster, etc.
- One script for end2end flow:
 - Data ingestion
 - Transformation
 - Analysis
 - Persist data
- Manages tasks for you
 - DAG analysis – Directed Acyclic Graph
 - Launch/Distribute/(Re)Launch...

Why Scala?

- A bit mind-bendy, but can code data transformations efficiently
- Schlepp functions around, not data
- Lazy eval
- Spark developed on Scala
 - GraphX only available in Scala
- Scala worksheets
 - Facilitates development in Scala for data scientists!

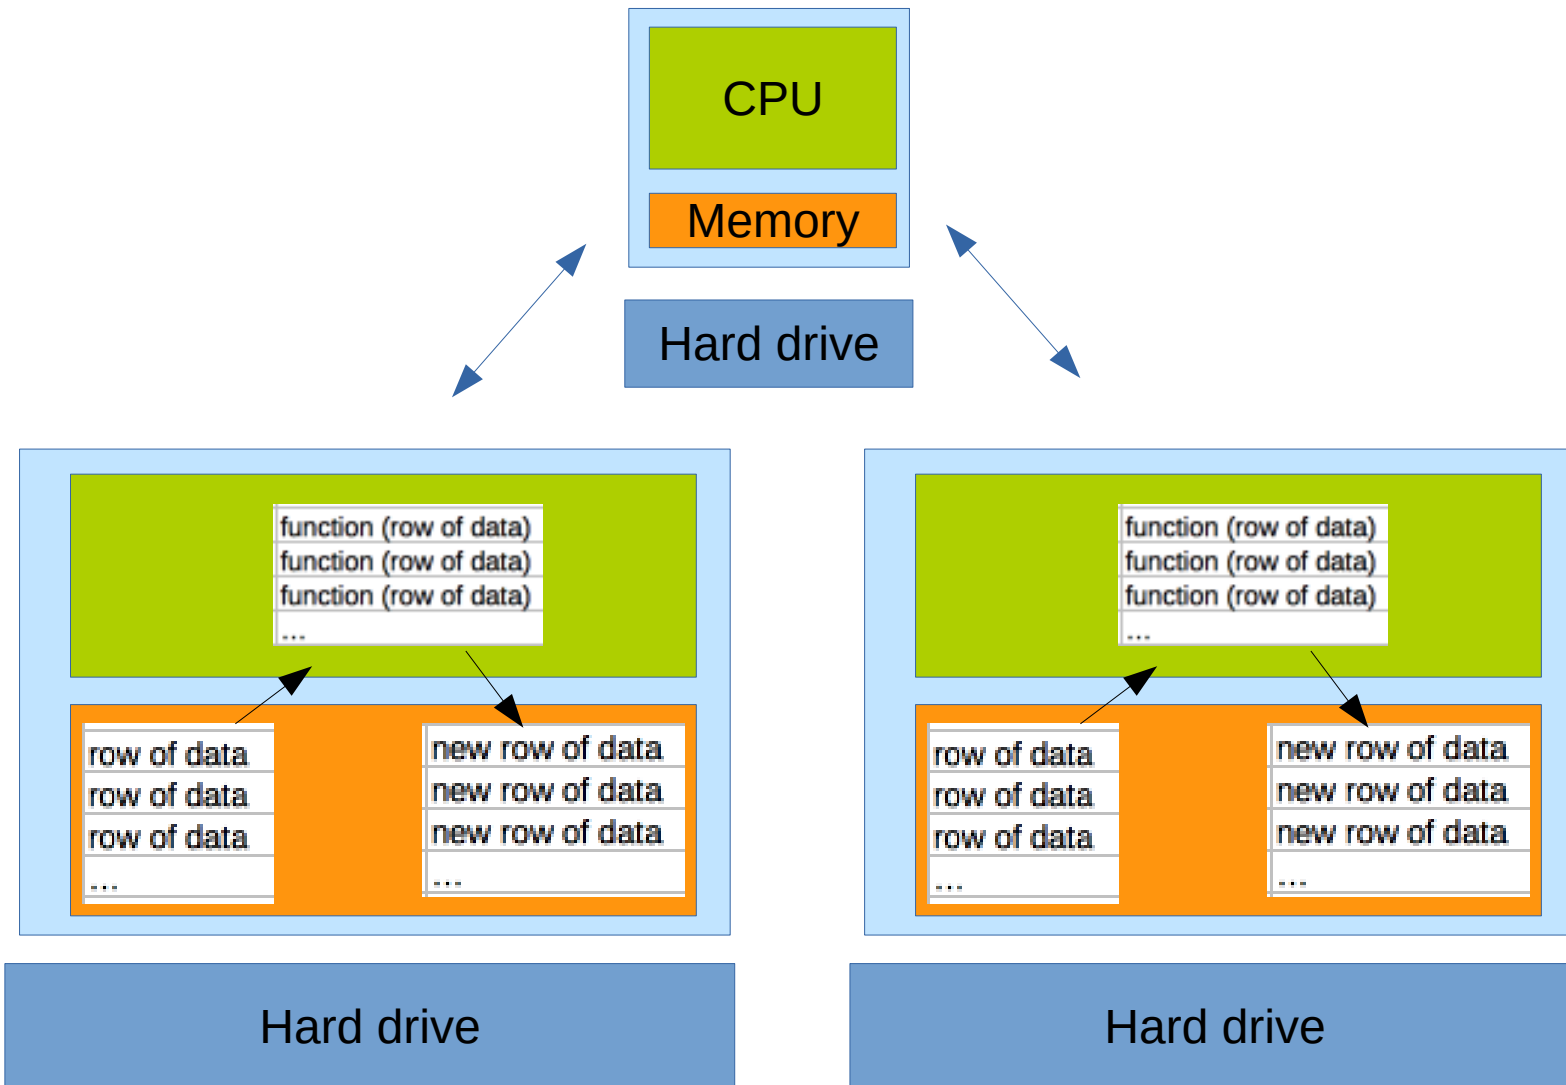
Airline Data for Demos

- Government data from airlines
 - http://www.transtats.bts.gov/DL_SelectFields.asp?Table_ID=236
- Airline, plane, flight info, etc
 - Clean up/reformatting required
 - Sample data available in data folder in git repo

Data Transformations

- What happens where?
 - Functions applied per row of data
 - Functions that aggregate rows
- Why is knowing this important?

Per Row



```

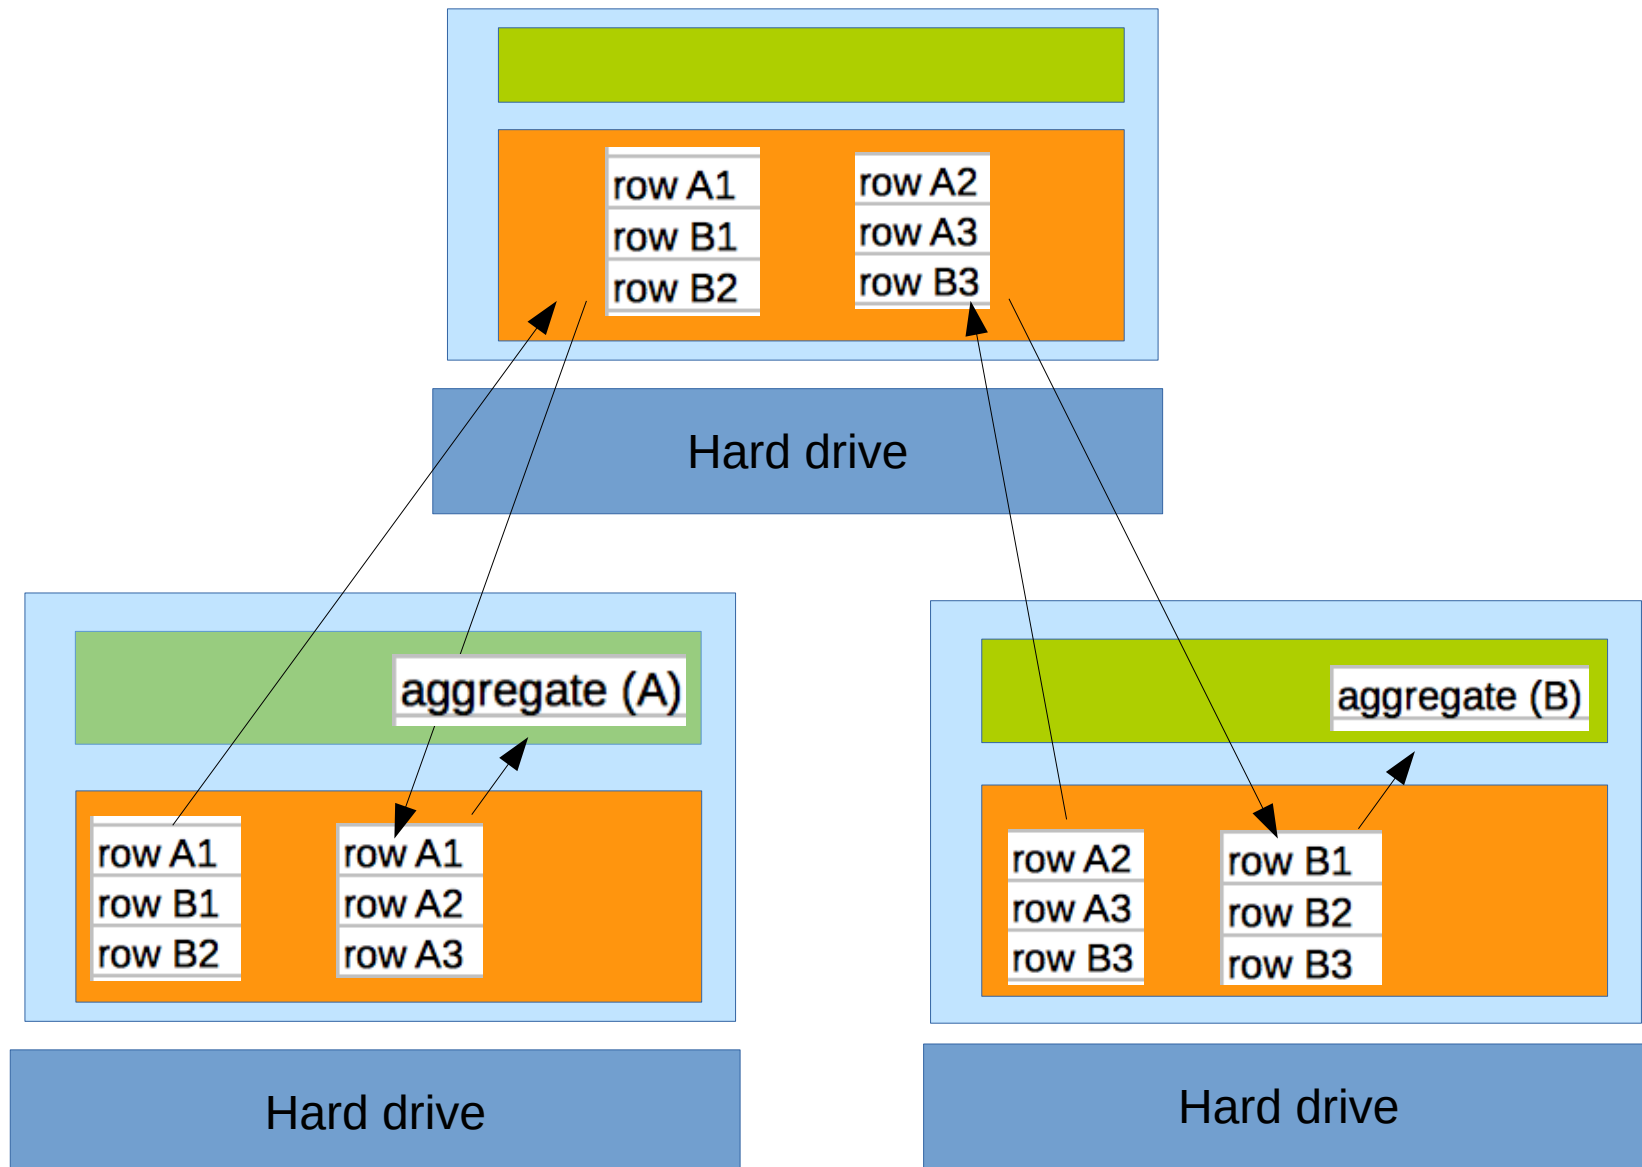
/* row transformations */
// This UDF (user defined function) takes an integer and converts it into a float that can be used for time calculations
val convertTime = udf((time: Integer) =>
  if (time > 100) time.toString().dropRight(2).toString().toFloat + time.toString().takeRight(2).toFloat/60
  else time.toString().toFloat/60
)
//> convertTime : org.apache.spark.sql.UserDefinedFunction = UserDefinedFunction(
//| on(<function1>,FloatType,List(IntegerType))

val aprilDF1 = aprilDF.withColumn("CRS_Dep_Time_f",convertTime(aprilDF("CRS_DEP_TIME")))
//> aprilDF1 : org.apache.spark.sql.DataFrame = [QUARTER: int, MONTH: int, DAY
//| _OF_WEEK: int, FL_DATE: string, AIRLINE_ID: int, CARRIER: string, TAIL_NUM:
//| string, FL_NUM: int, ORIGIN_AIRPORT_ID: int, ORIGIN: string, DEST_AIRPORT_
//| ID: int, DEST: string, CRS_DEP_TIME: int, DEP_TIME: int, TAXI_OUT: double,
//| WHEELS_OFF: int, WHEELS_ON: int, TAXI_IN: double, CRS_ARR_TIME: int, ARR_TI
//| ME: int, DISTANCE: double, : string, CRS_Dep_Time_f: float]

aprilDF1.select("CRS_DEP_TIME","CRS_Dep_Time_f").take(5)
//> res0: Array[org.apache.spark.sql.Row] = Array([900,9.0], [900,9.0], [900,9.
//| 0], [900,9.0], [900,9.0])

```

The Shuffle



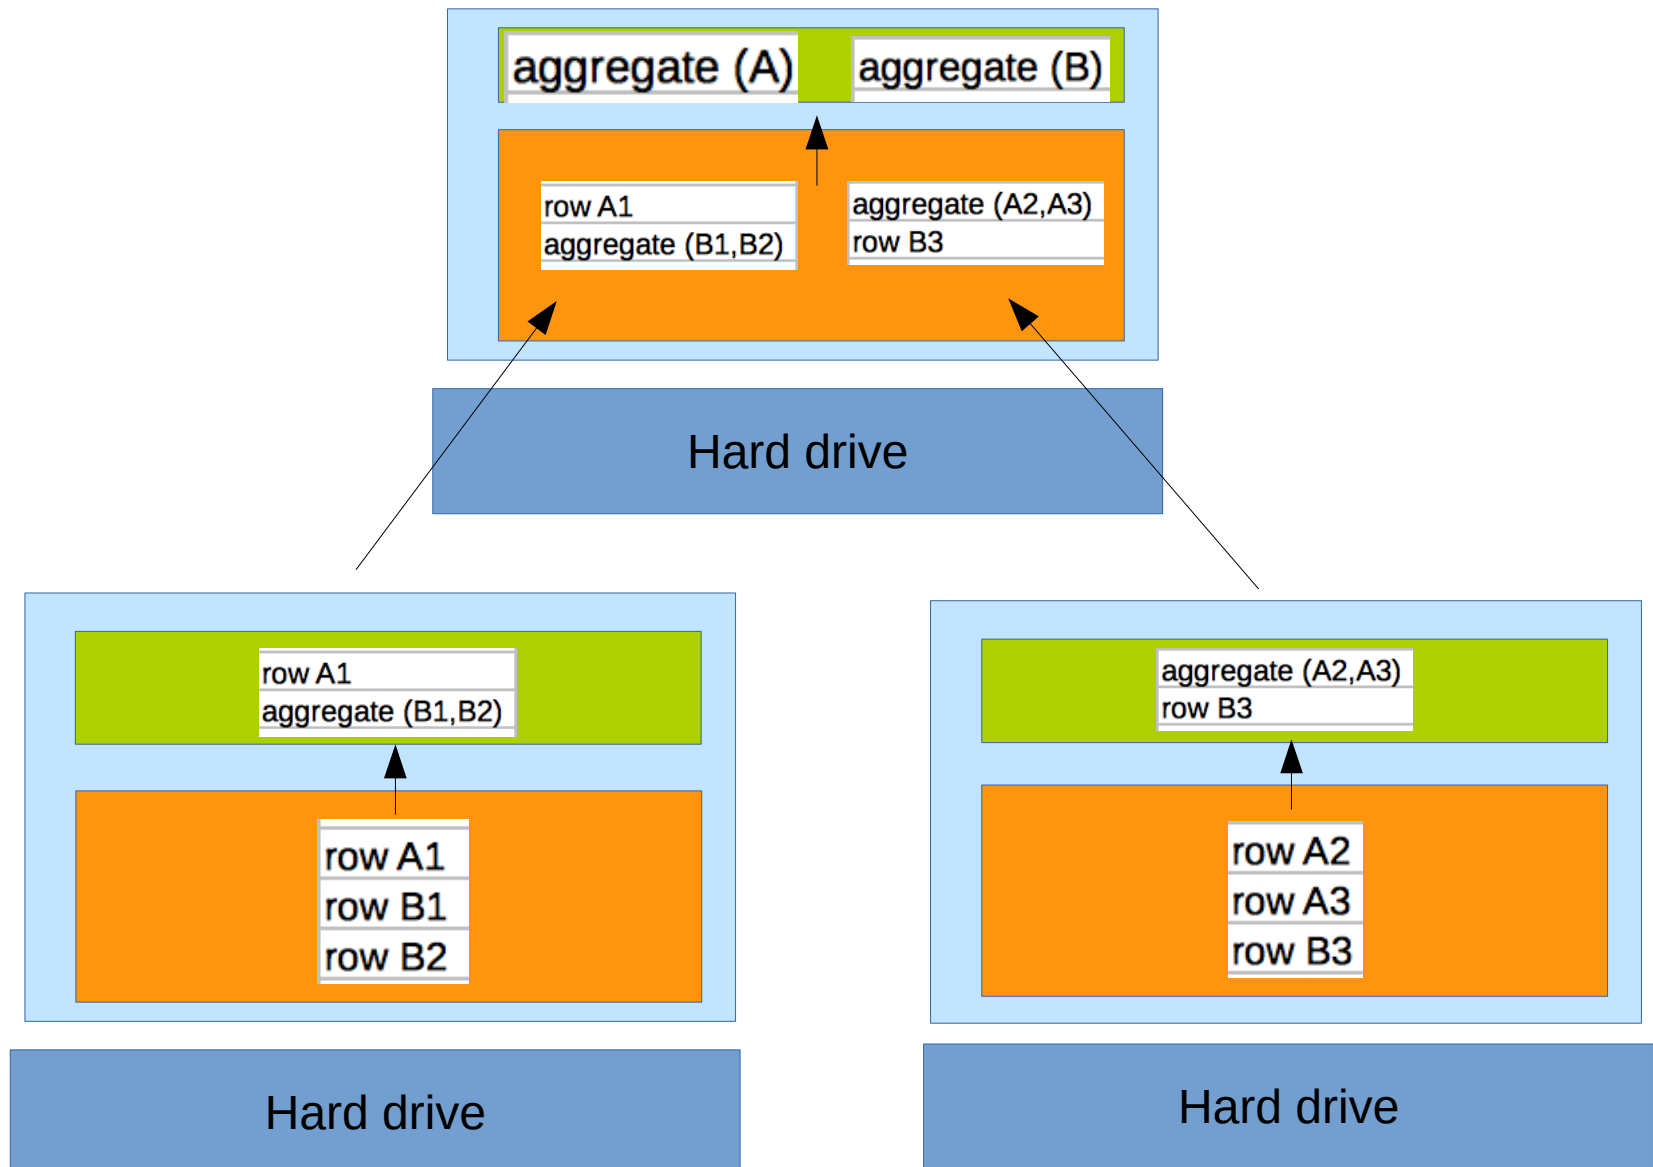
```

/* demo of groupBy on PairRDDs */
// create PairRDD of AIRLINE_ID and count (1)
val aprilPairRDD = aprilDF.select("AIRLINE_ID").map(id => (id,1))
//> aprilPairRDD : org.apache.spark.rdd.RDD[(org.apache.spark.sql.Row, Int)] =
//| MapPartitionsRDD[30] at map at Spark4DS.developDemo.scala:50

// do a groupBy on PairRDD (causes full shuffle)
aprilPairRDD.groupByKey().map(t => (t._1, t._2.sum)).collect()
//>
[Stage 7:>                                (0 + 2
//| ) / 2]
[Stage 7:=====>
//| (1 + 1) / 2]
[Stage 8:>
//|      (0 + 2) / 2]
[Stage 8:=====>
//|      (1 + 1) / 2]
//|
res1: Array[(org.apache.spark.sql.Row, Int)]
//| = Array(([19930],13974), ([20304],49329), ([21171],4915), ([20355],32496),
//| ([19805],44770), ([20416],9496), ([19790],72170), ([20398],25695), ([20436
//| ],7148), ([20366],49296), ([19977],41342), ([19393],106407), ([20409],22020
//| ), ([19690],6093))

```

(Semi) Aggregate, then Shuffle



```

/* demo of reduceBy on same PairRDD */      //>
// reduceBy does an aggregation first within each partition and then does a partial shuffle
aprilPairRDD.reduceByKey(_ + _).collect()
[Stage 9:>                                (0 + 2

//| ) / 2]

//|
res2: Array[(org.apache.spark.sql.Row, Int)] = Array((19930],
//| 13974), ([20304],49329), ([21171],4915), ([20355],32496), ([19805],44770),
//| ([20416],9496), ([19790],72170), ([20398],25695), ([20436],7148), ([20366],
//| 49296), ([19977],41342), ([19393],106407), ([20409],22020), ([19690],6093))
//|

```

Spark DF GroupBy

```

aprilDF1.groupBy("AIRLINE_ID").count().show() //>
[Stage 3:>                                     (0 + 2)

//| ) / 2]

//|

[Stage 6:=====>
//| (125 + 4) / 199]
[Stage 6:=====
//| => (173 + 4) / 199]

//|

+-----+

//| |AIRLINE_ID| count|
//| +-----+-----+
//| |    20436|   7148|
//| |    19690|   6093|
//| |    20304|  49329|
//| |    19930|  13974|
//| |    20355|  32496|
//| |    20366|  49296|
//| |    21171|   4915|
//| |    19977|  41342|
//| |    19790|  72170|
//| |    19393|106407|
//| |    20398|  25695|
//| |    19805|  44770|
//| |    20409|  22020|
//| |    20416|   9496|
//| +-----+-----+
//|

```


Summary: Spark for Data Scientists

- Think
 - Scala Spark
 - Hardware AND software!
- Next steps for airline data analysis
 - Are there factors that indicate a propensity for being late?
 - What is a good metric for lateness?
 - What other derived variables can you create (in a performance optimized way!) to help your analysis?
 - Can you identify airline hubs using the data?
 - GraphX and connected components?

Appendix

- Build Spark with Scala 2.11
- Eclipse with Scala
- Scala Worksheets
 - New worksheet
 - Worksheet settings
- Build using SBT

Build Spark with Scala 2.11

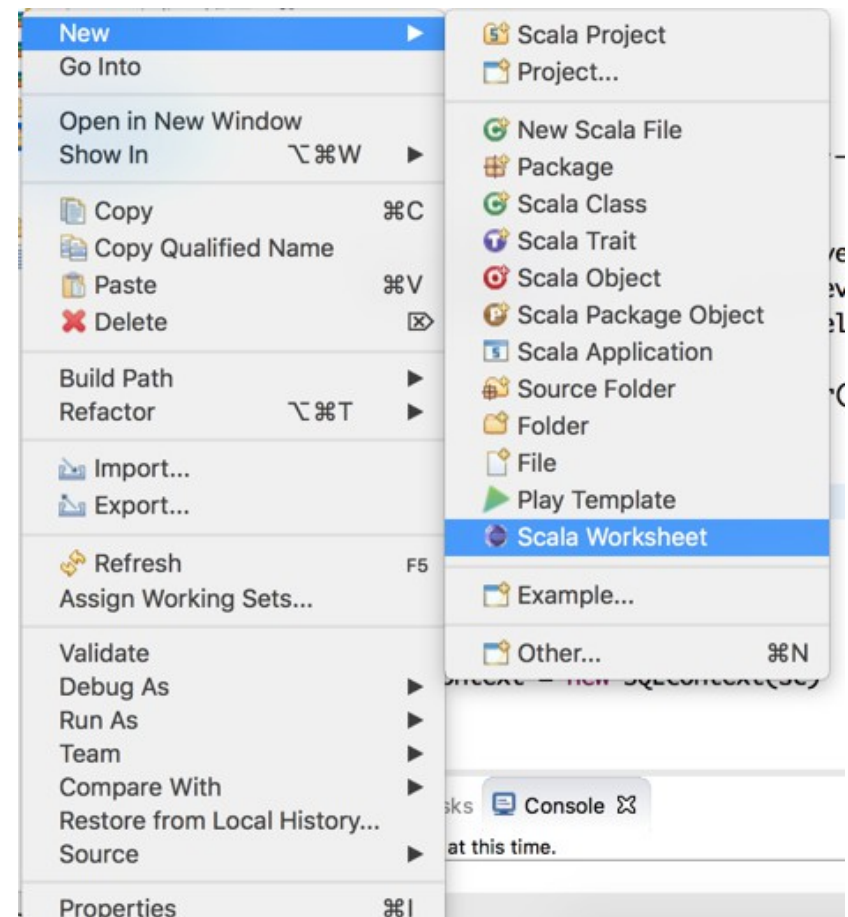
- Instructions developed for MacBook Pro running OS X El Capitan 10.11.2
- Download source file from <http://spark.apache.org/downloads.html>
 - Should get a file called spark-version.tgz
- Move file to /usr/local/
 - `mv /Users/username/Documents/spark-1.6.0.tgz /usr/local/`
- Unpack
 - `tar -xf spark-1.6.0.tgz`
- Delete zipped file
 - `rm spark-1.6.0.tgz`
- Change to spark directory
 - `cd spark-1.6.0`
- Explicitly set Scala version (May or may not need to do this step)
 - `./dev/change-scala-version.sh 2.11`
- Build using maven (also adding in Yarn, Hadoop (2.6.0), and Hive)
 - `mvn -Pyarn -Phadoop-2.6 -Dhadoop.version=2.6.0 -Phive -Phive-thriftserver -Dscala-2.11 -DskipTests clean package`
- Run Spark from this folder (or set alias to this folder)
 - `bin/run-example SparkPi 10`
 - `bin/spark-shell`
 - `bin/spark-submit`

Eclipse with Scala

- Download from here
 - <http://scala-ide.org/download/sdk.html>
 - Need Scala 2.11.7 version for worksheets to work with Spark!
 - If you already have the Java version of Eclipse & want to use the Scala plugin...
 - Good luck. Scala Worksheets & some other functionality was really buggy the last time I tried it (2014), but maybe it's grown up since then... I just use two separate Eclipse apps (one for Scala & one for Java)
- IDE installation should include Scala & SBT
 - See notes on Building using SBT page for working with Eclipse from SBT

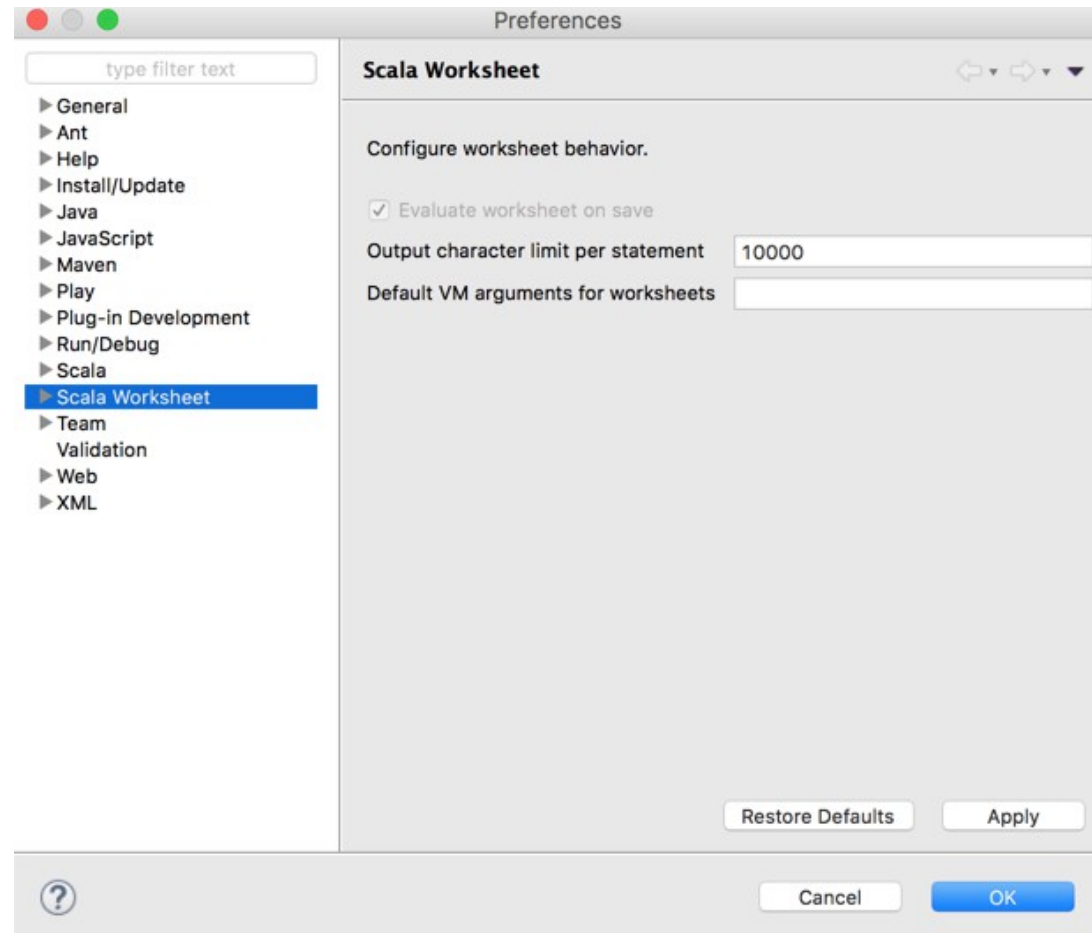
Create New Scala Worksheet

- Create new Scala project
- Create new Scala worksheet in project



Scala Worksheet Settings

- Scala IDE/Preferences
 - Evaluate worksheet on save
 - Output character limit



Build using SBT

- Need build.sbt in your project folder
 - See example build.sbt in git repo
- Need plugins in sbt plugin folder (for Eclipse & assembly)
 - .sbt/0.13/plugins/plugins.sbt
 - See example plugins.sbt in git repo
- Creating new projects
 - Directly in Eclipse
 - Manual creation and/or cloning git repos
 - cd to project folder
 - Create files needed by eclipse by typing following command
 - sbt eclipse
- Building executable jar
 - cd to project folder
 - sbt
 - clean
 - (compile)
 - assembly
 - Executable jar will be in project/target/Scala folder